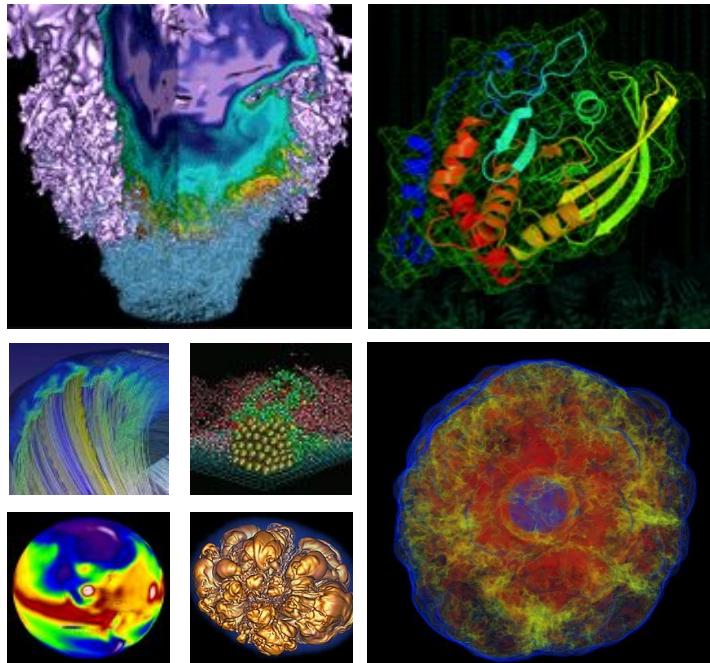


# Optimizing Python-Based Spectroscopic Data Processing on NERSC Supercomputers



Laurie Stephey  
SciPy 2019, July 10 2019

Rollin Thomas, NERSC and Stephen Bailey, DESI  
Lawrence Berkeley National Laboratory



Office of  
Science



# Overview of this talk

---



- DESI Python code now 5-7x faster!
- How?
  - Profiled to find hotspots
  - Major restructures
  - JIT compiled with Numba
- MPI vs Dask for parallelization (chose MPI)
- Current and future work: GPUs

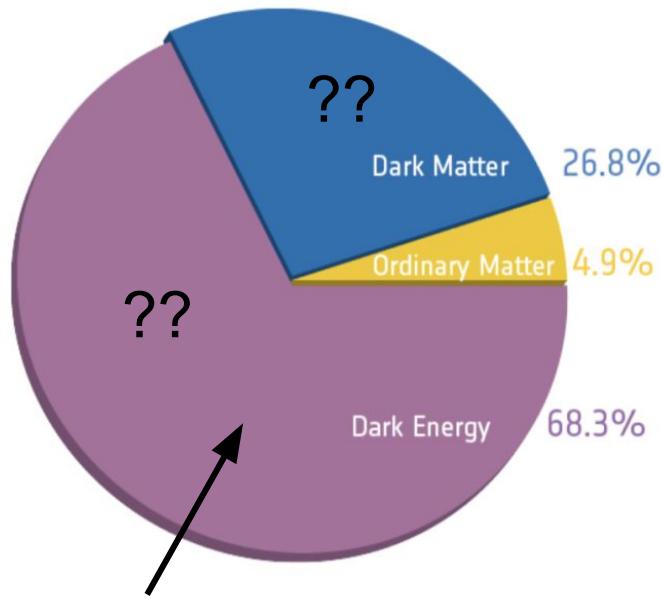
# What is DESI?



## Dark Energy Spectroscopic Instrument



The Mayall telescope on Kitt Peak, AZ, where the DESI instrument is currently being installed



We currently believe the universe is almost 70 percent dark energy, but we don't have a good understanding of what dark energy actually is!

# A 3D map of the universe!



- Over 5 years, DESI will make the most detailed 3D map of the universe ever!
- They will use this map to better understand what dark energy is and its role throughout the evolution of the universe
- I am not a cosmologist, so for more (and better!) information, check out  
<https://www.desi.lbl.gov/>

# What is NERSC?



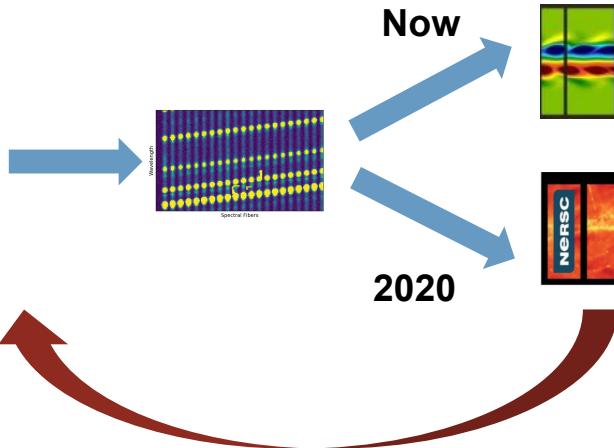
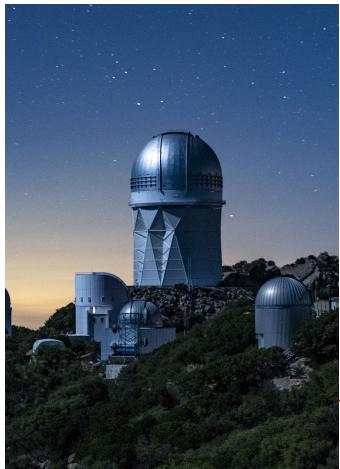
- Workhorse computing facility for the Office of Science in the US Department of Energy
- Located at Lawrence Berkeley Lab in Berkeley, California
- 7000 users
- Cori: current system (Intel Haswell + KNL)
- Perlmutter: future system (AMD CPUs + NVIDIA GPUs)



# Why worry about speed?



DESI will take hundreds of images every night for 5 years



Will send images to NERSC to be quickly processed



Processed data used to make decisions for next night of observing

This will require many millions of CPU hours! Can we make it more efficient?

# Project goals and constraints



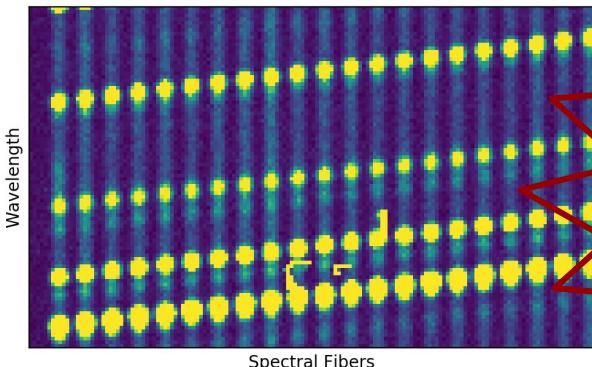
- DESI commissioning starts this fall
- Need fast turnaround for survey operations and scientific results
- DESI prioritizes quick, clear, maintainable code → Python! (This audience gets it! 😊)
- Asked that we don't rewrite code in C or Cython (harder to write, read, and maintain)
- Our mission: make spectral extraction code faster!

# What is spectral extraction?



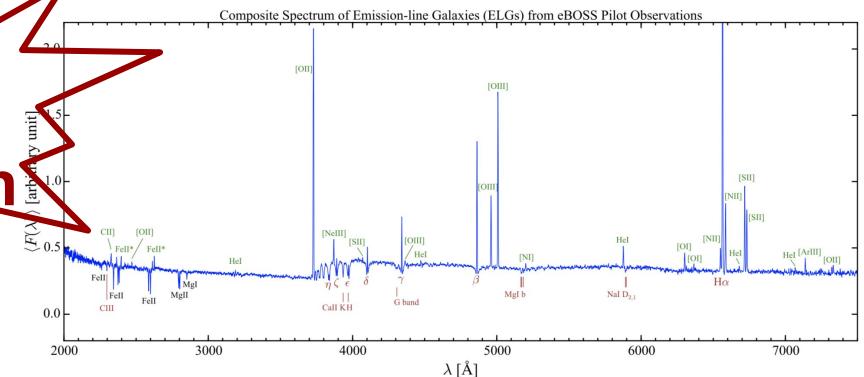
- Take raw image data and turn it into usable spectra
  - “Spectroperfectionism” 2D fitting (expensive!)
  - Eigenvalue decomposition
  - Evaluating special functions
  - Bookkeeping...

## Raw data on CCD



# Spectral extraction

# Goal: High quality output spectra



# Where do we begin?



desihub / specter

Unwatch 47 ★ Star 4 Fork 3

Code

Issues 15

Pull requests 0

Projects 0

Wiki

Insights

A toolkit for simulating multi-object spectrographs

551 commits

6 branches

25 releases

5 contributors

View license

Branch: master ▾

New pull request

Create new file

Upload files

Find file

Clone or download ▾

sbailey updated changes.rst for PR #71 faster pgh

Latest commit a70f99e 10 days ago

bin	add unicode_literals	2 years ago
dev	add bright star on border case to divide-and-conquer notebook	a year ago
doc	updated changes.rst for PR #71 faster pgh	10 days ago
etc	update module file	2 years ago
py/specter	Merge pull request #71 from desihub/fix_pgh	10 days ago
.coveragerc	adding top-level files	3 years ago
.gitignore	Specter speedup (#56)	10 months ago
.travis.yml	.travis.yml debugging	5 months ago
LICENSE.rst	fix license file	a year ago
MANIFEST.in	update MANIFEST.in	2 years ago
README.rst	remove incorrect versioning instructions	2 years ago
setup.py	update release notes and version for specter/0.8.2	10 months ago

Can you find  
the slow part?



# Profiling tools: cProfile

---



- **Easy! Built-in, no modifications to your code:**
- **Example:**

```
python -m cProfile -o output.prof path/to/your/script arg1 arg2
```

- **Writes human readable files, but nicer to visualize**

# Profiling tools: cProfile



cProfile results visualized with [Snakeviz](#) and [gprof2dot](#)

Style: Sunburst

Depth: 10

Cutoff: 1 > 1000

Name:  
\_xypix  
Cumulative Time:  
78.2 s (51.30 %)

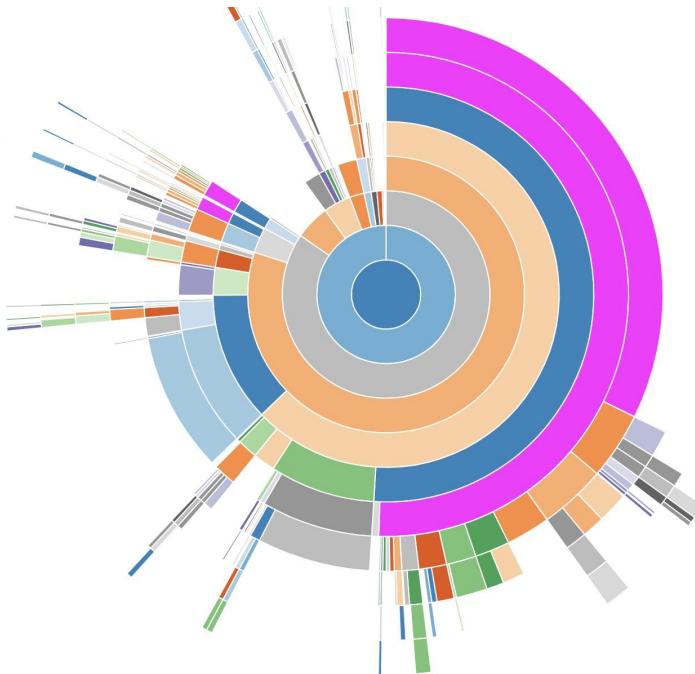
File:  
gausshermite.py

Line:

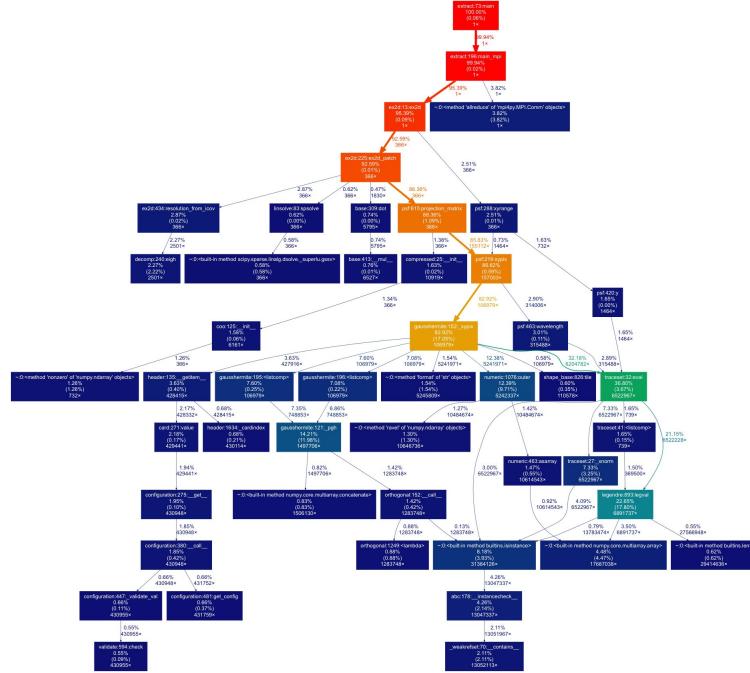
124

Directory:

/global/cscratch1/sd/st  
ephey/git\_repo/specter/  
py/specter/psf/



Snakeviz



gprof2dot

# Profiling tools: [line profiler](#)

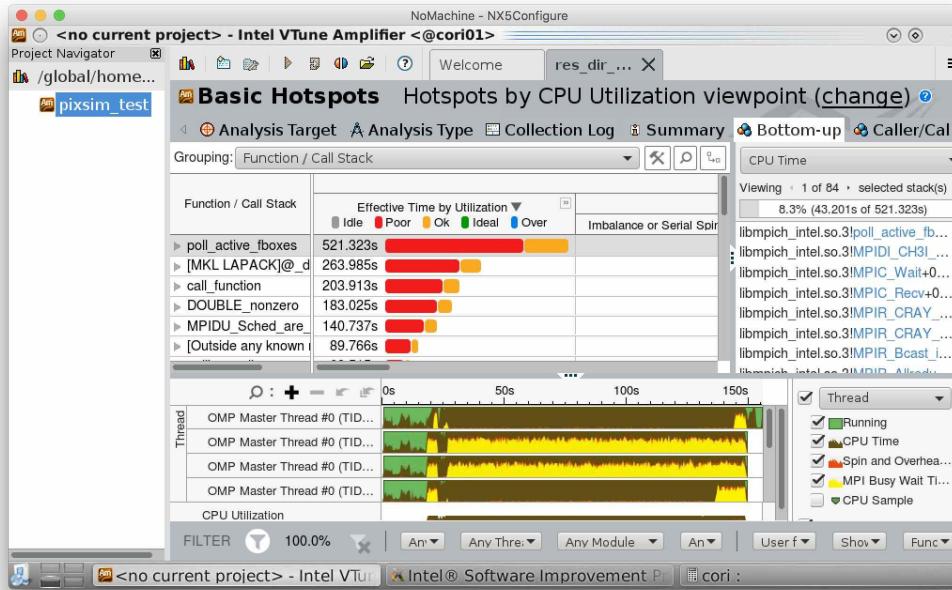


- Manually decorate your functions of interest
  - Get line-by-line profiling information

# Profiling tools: Intel Vtune



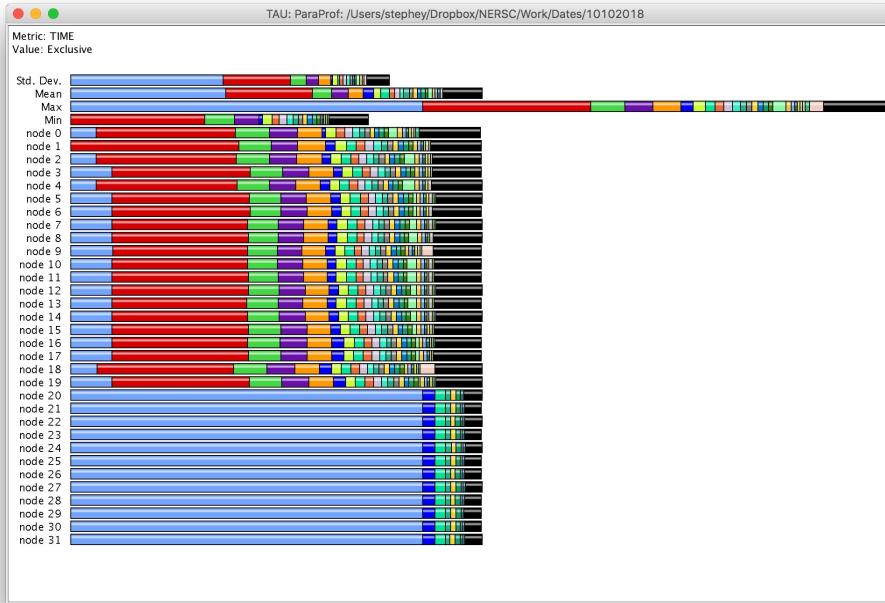
- Vtune has some support for Python
- Tricky to use, collect, and display the right data



# Profiling tools: Tau (U Oregon)



- For best results, custom build for your application
- Excellent and interactive visualizations
- Great for MPI applications



# Our final profiling strategy

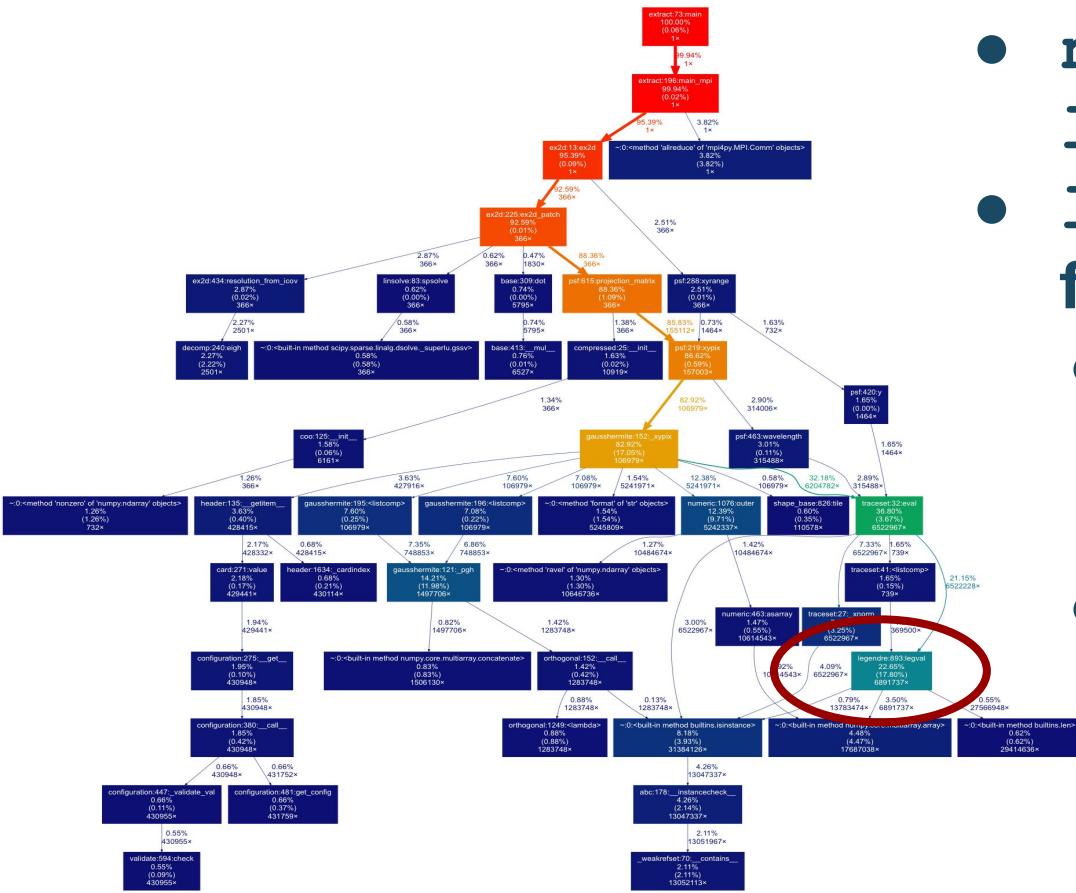


- Start by profiling the code with `cProfile`, visualize
- Identify a hotspot
- Examine the function with `line_profiler` to get more information
- Decide on a strategy for the hotspot:
  - Speed up function by JIT-compiling (easier)
  - Restructure the code (harder)
- Benchmark and reprofile the code, see how well it worked (or didn't)
- Rinse and repeat, use “big” tools only when required

# Hotspot: `legval`



- `numpy.polynomial.legendre.legval`
  - `legval` was a problem for two reasons:
    - Expensive function (in Python rather than C)
    - Called too many times (often with scalar values!)



# Strategy #1 JIT compile



- **Numba is a JIT compiler for Python (LLVM under the hood)**
- Get the function into a form Numba will accept (can be tricky)
  - Remove type checking, other unsupported functionality
  - Remove all SciPy (not currently supported) 😞
  - Size and type of arrays must be known at compile time
- Function is now faster!



Office of  
Science



# JIT compile legval



## Before

```
def legval(x, c, tensor=True):
    c = np.array(c, ndmin=1, copy=0)
    if c.dtype.char in '?bBhHiIlQqP':
        c = c.astype(np.double)
    if isinstance(x, (tuple, list)):
        x = np.asarray(x)
    if isinstance(x, np.ndarray) and tensor:
        c = c.reshape(c.shape + (1,)*x.ndim)

    if len(c) == 1:
        c0 = c[0]
        c1 = 0
    elif len(c) == 2:
        c0 = c[0]
        c1 = c[1]
    else:
        nd = len(c)
        c0 = c[-2]
        c1 = c[-1]
        for i in range(3, len(c) + 1):
            tmp = c0
            nd = nd - 1
            c0 = c[-i] - (c1*(nd - 1))/nd
            c1 = tmp + (c1*x*(2*nd - 1))/nd
    return c0 + c1*x
```

Removed  
this, Numba  
won't allow  
type checking

Removed these  
cases since we  
know  $\text{len}(c) > 2$

## After

```
@numba.jit(nopython=True, cache=False)
def legval_numba(x, c):
    nd=len(c)
    ndd=nd
    xlen = x.size
    c0=c[-2]*np.ones(xlen)
    c1=c[-1]*np.ones(xlen)
    for i in range(3, ndd + 1):
        tmp = c0
        nd = nd - 1
        nd_inv = 1/nd
        c0 = c[-i] - (c1*(nd - 1))*nd_inv
        c1 = tmp + (c1*x*(2*nd - 1))*nd_inv
    return c0 + c1*x
```

Add Numba decorator

Added extra  
info so the  
arrays don't  
change size

Some changes → 16x speedup on KNL, 5x  
speedup on Haswell for this function!



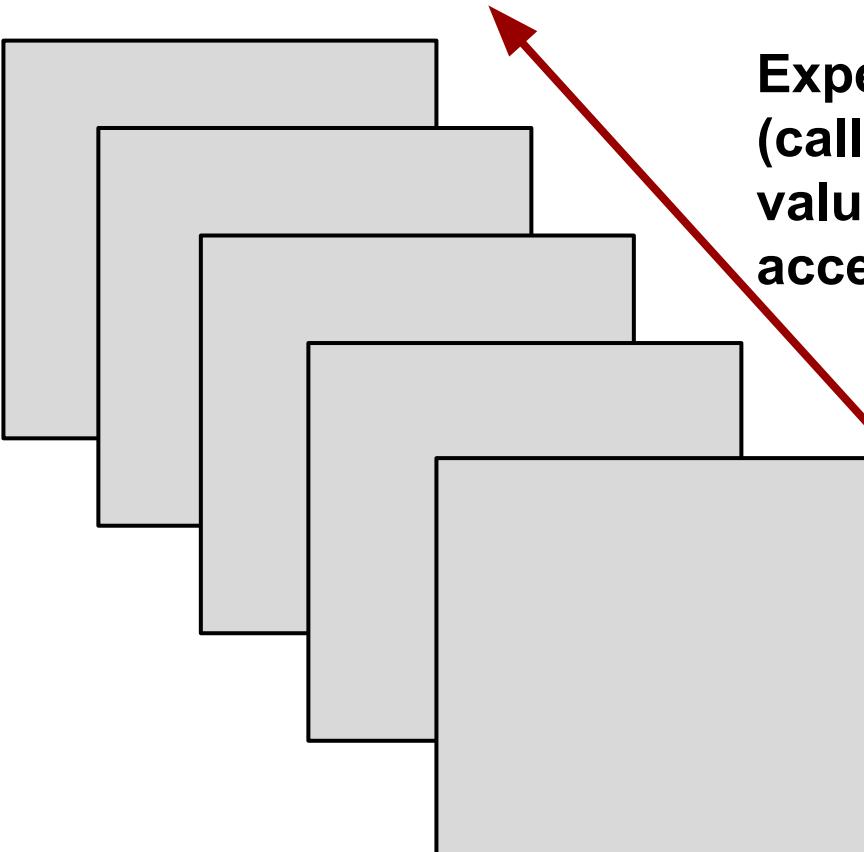
Office of  
Science



# Strategy #2: Restructure



High level



Low level

Expensive function  
(called fewer times,  
values stored and  
accessed later)

Expensive  
function (called  
many times)

# Restructure to mitigate legval



## Before

```
#- Extract GH degree and sigma coefficients for convenience
sigx1 = self.coeff['GHSIGX'].eval(ispec, wavelength)
sigy1 = self.coeff['GHSIGY'].eval(ispec, wavelength)

#- Background tail image
tailxsca = self.coeff['TAILXSCA'].eval(ispec, wavelength)
taileysca = self.coeff['TAILYSCA'].eval(ispec, wavelength)
tailamp = self.coeff['TAILAMP'].eval(ispec, wavelength)
tailcore = self.coeff['TAILCORE'].eval(ispec, wavelength)
tailinde = self.coeff['TAILINDE'].eval(ispec, wavelength)
```



**Calling eval (which calls legval) again and again for scalars**



## After

```
#- Extract GH degree and sigma coefficients for convenience
sigx1 = self.legval_dict['GHSIGX'][ispec_cache, iwave_cache]
sigy1 = self.legval_dict['GHSIGY'][ispec_cache, iwave_cache]

#- Background tail image
tailxsca = self.legval_dict['TAILXSCA'][ispec_cache, iwave_cache]
taileysca = self.legval_dict['TAILYSCA'][ispec_cache, iwave_cache]
tailamp = self.legval_dict['TAILAMP'][ispec_cache, iwave_cache]
tailcore = self.legval_dict['TAILCORE'][ispec_cache, iwave_cache]
tailinde = self.legval_dict['TAILINDE'][ispec_cache, iwave_cache]
```



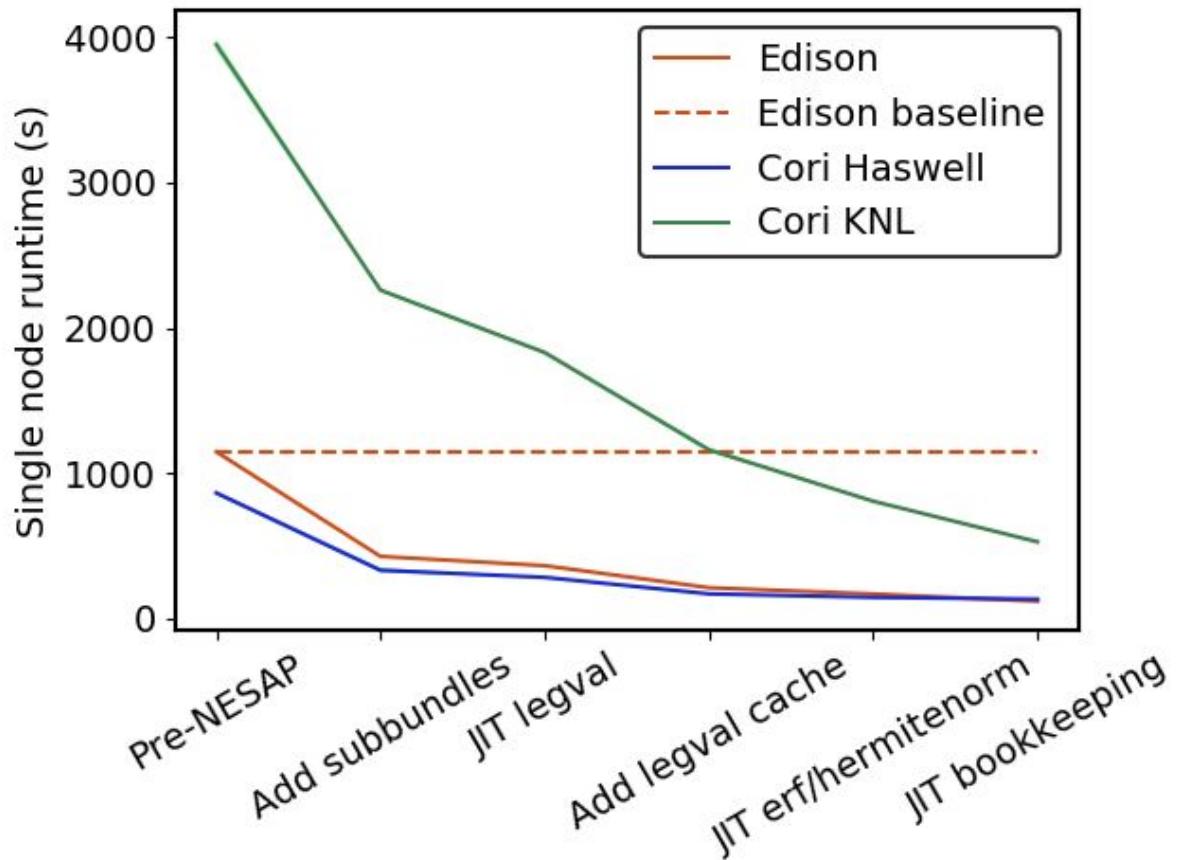
**Just looking up legval values we have already computed!**



# Is it faster? Yes!



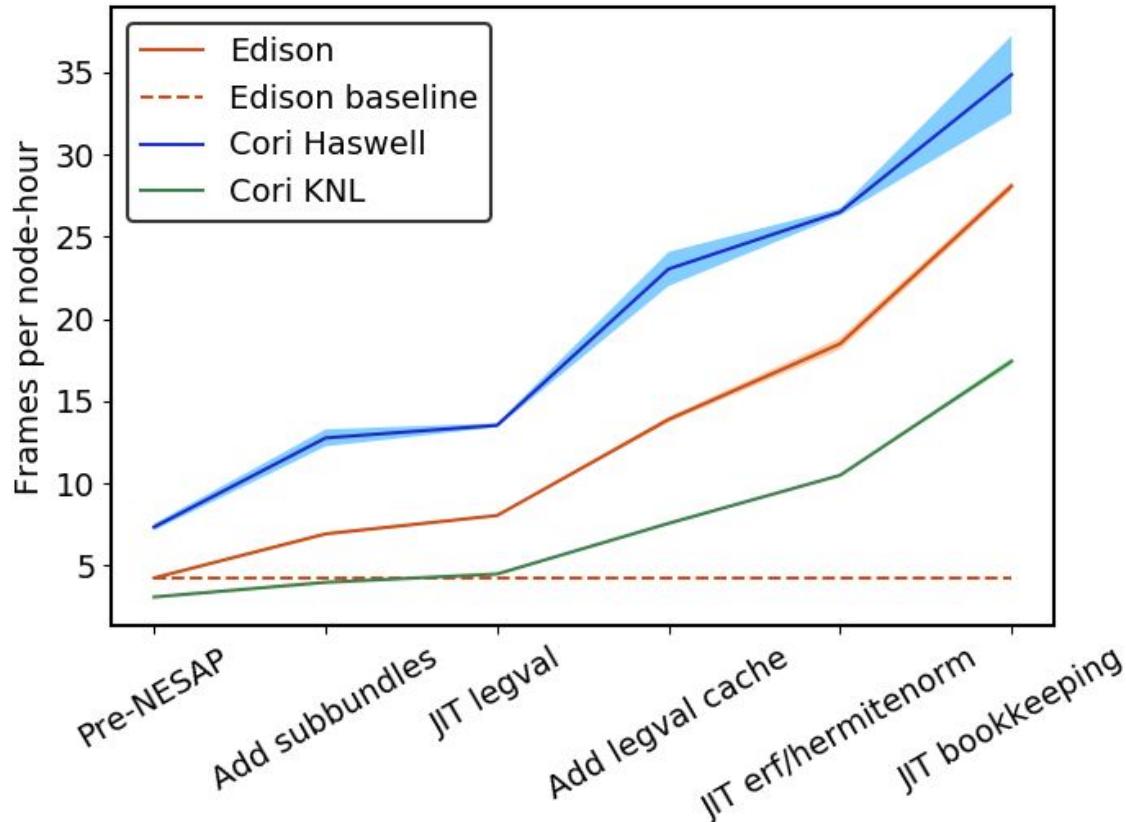
**DESI spectral extraction now  
7-10x faster for a  
single node**



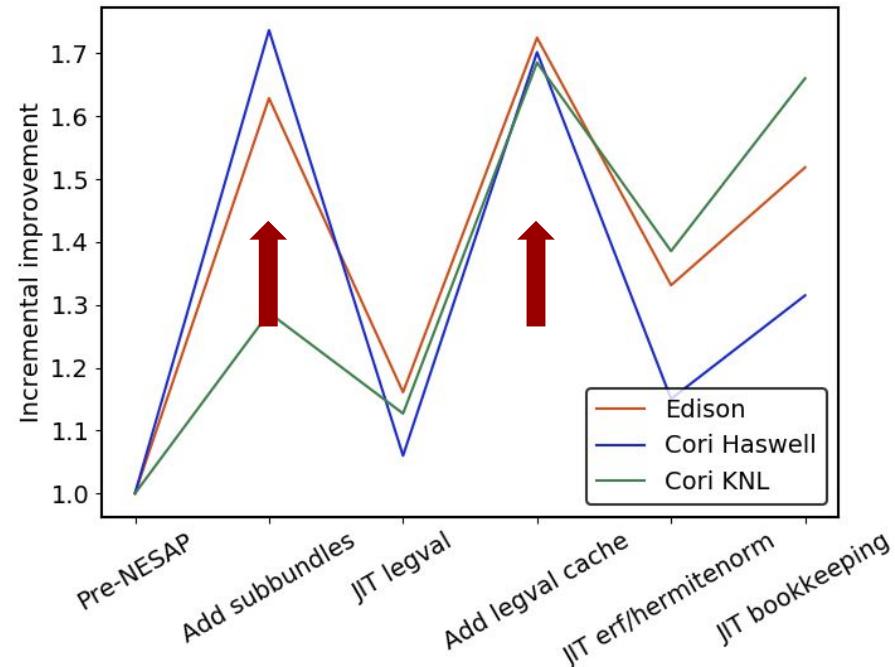
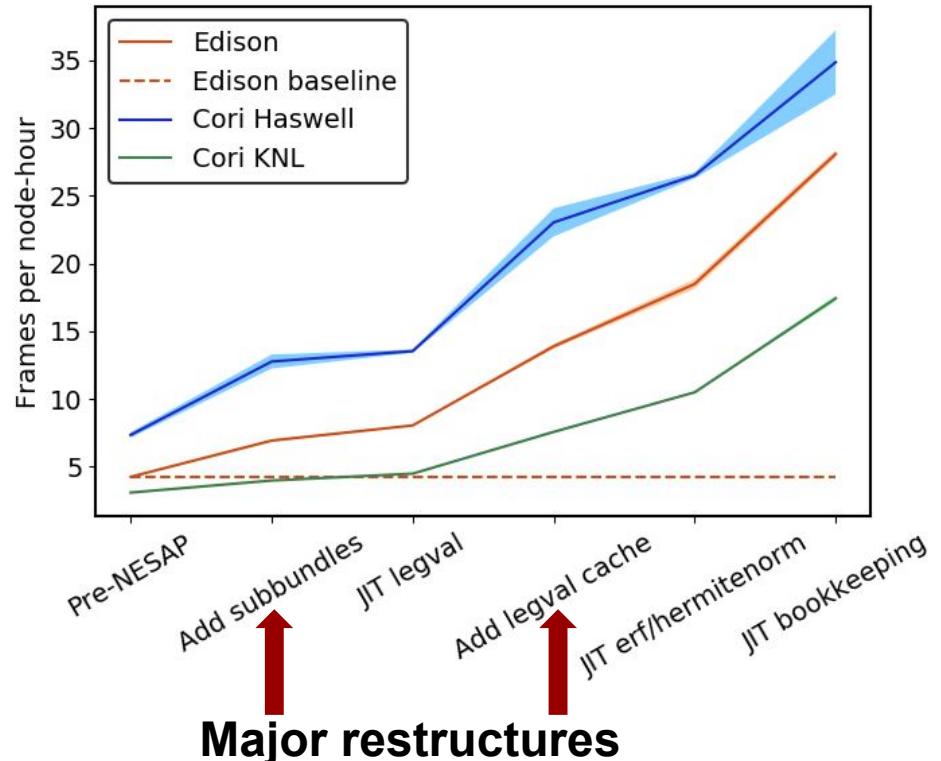
# 5-7x more throughput!



- For DESI, throughput is more meaningful than sheer speed
- How many frames can they process per CPU-hour?
- **5-7x more throughput**
- Before: 33 million CPU hours
- After: 6.5 million CPU hours

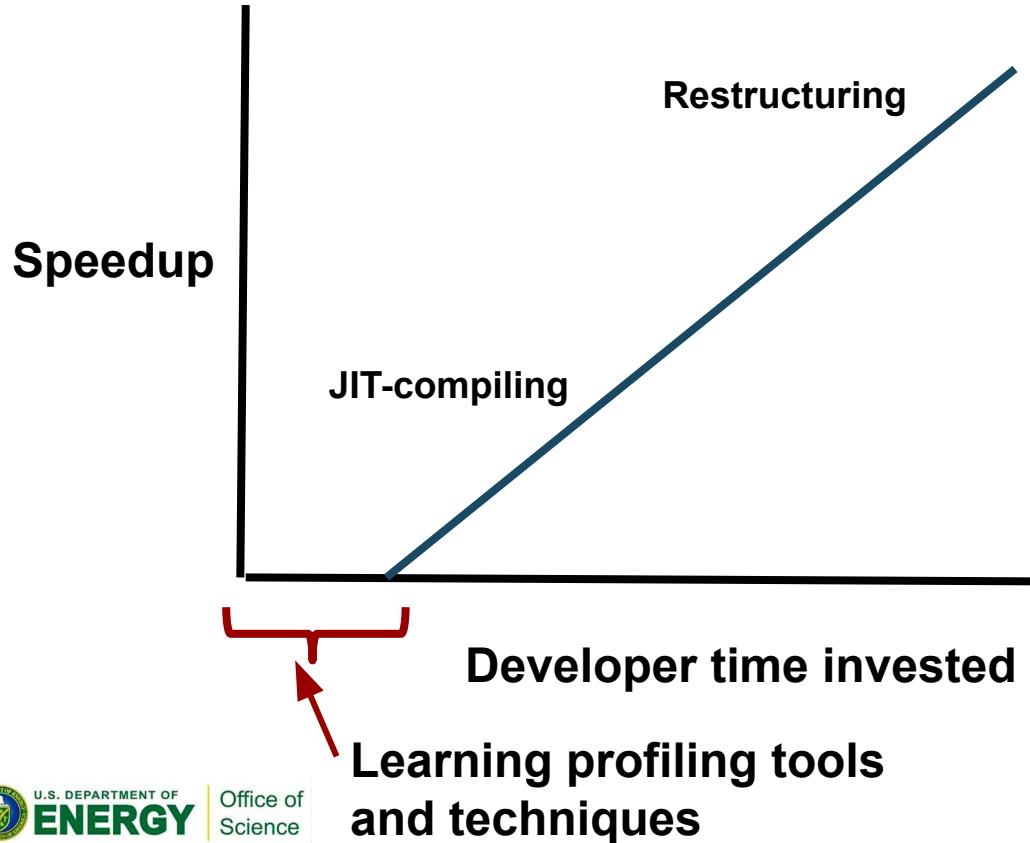


# Restructures → more speedup



**Restructures resulted in more incremental speedup than JIT-based optimizations**

# Our experience: time in $\infty$ speedup out

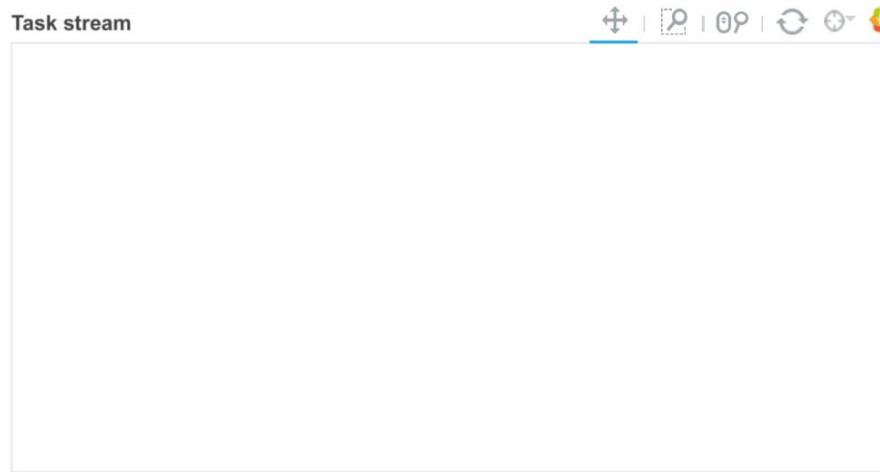


- Ultimate speedup is proportional to time invested (i.e. speedup is not a step function)
- \*\*\*No actual data here, just our experience!
- tl;dr: If you don't have much time, you can still obtain modest improvements

# What about Dask?



- We like Dask and wondered if the DESI mpi4py could be replaced with Dask
- Short answer: no (too late!)



GIF reproduced from: <https://distributed.dask.org/en/latest/web.html>

# “Change is the only constant in life”



CPUs



- Heraclitus,  
Greek philosopher

CPUs + GPUs

- Coming in 2020
- Time to start preparing applications now

# DESI GPU open questions

---



- How can DESI best leverage the GPUs?
- Should we focus on speeding up functions like `legval` in a bottom-up approach like we did for CPUs?
- Should we take a top-down approach and rethink our core algorithms? (Probably)
- What is the best option for Python? CuPy + Numba?
- This is our current/future work

# Summary

---



- We were able to help DESI improve their throughput 5-7x without rewriting in C!
- 33 million cpu hours → 6.5 million cpu hours
- Used two strategies to obtain speedup:
  - JIT-compiling with Numba (easy, but less overall improvement)
  - Restructuring of code (hard, but more overall improvement)
- Future work: port DESI code to GPUs

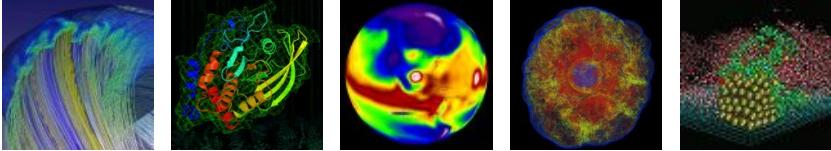
# Want to know more?

---

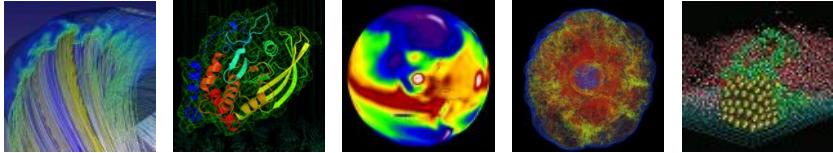


- Check out our documentation for profiling your Python code:  
<https://docs.nersc.gov/programming/high-level-environments/python/profiling-python/>
- Check out the DESI code on github:  
<https://github.com/desihub>
- Check out our SciPy 2019 proceedings paper:  
[http://conference.scipy.org/proceedings/scipy2019/laurie\\_stephey.html](http://conference.scipy.org/proceedings/scipy2019/laurie_stephey.html)

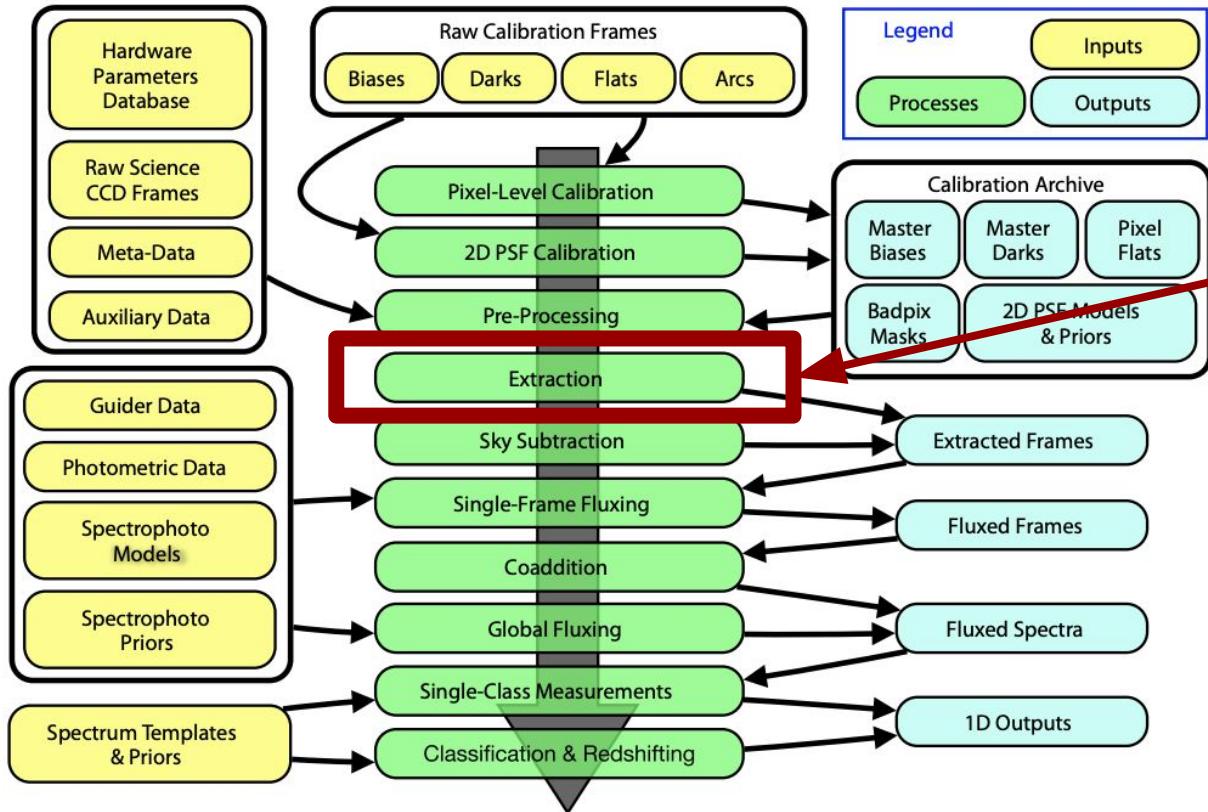
# Thank you!



# Extra slides



# The DESI pipeline

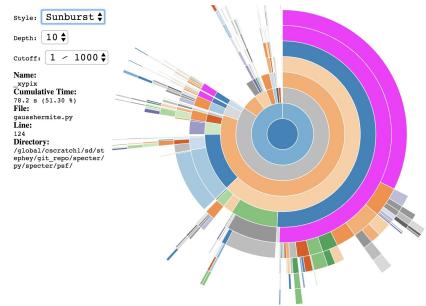


- **The spectral extraction part of the pipeline was known to be a major bottleneck**
- **Our mission: speed it up!**

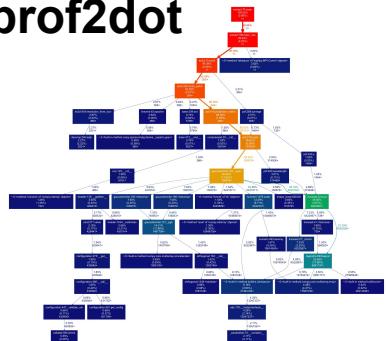
# Profiling workflow: easy to hard



# cProfile (built into Python) Snakeviz



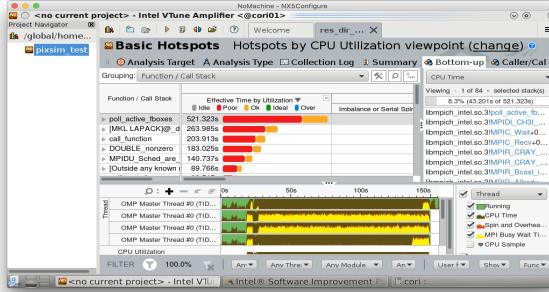
# gprof2dot



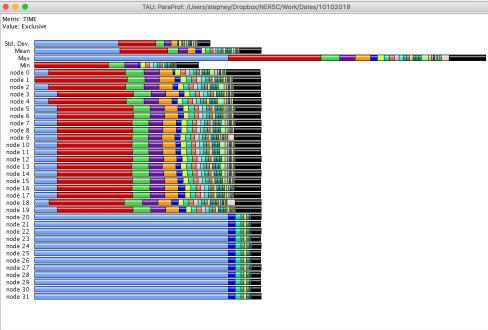
# Easy



# line\_profiler



Tau



# Hard



# DASK wasn't the right fit for DESI



- DESI uses mpi4py for MPI parallelism
- If DESI were using the concurrent futures API (MPI comm.spawn), it would have been conceptually similar to DASK and easier to transition
- Alas, this is not yet supported on NERSC Cori (due to Slurm compatibility)
- Ultimately we decided that changing to DASK was impractical for DESI
- Starting a new project with DASK in mind would be better than trying to adapt an established project



Office of  
Science



# CuPy and Numba



- We have tested CuPy and Numba for GPU Python
- CuPy:
  - Almost a drop-in replacement for NumPy
  - Easy to use
  - Not all NumPy/SciPy is implemented
- Numba:
  - Harder to use than its CPU cousin
  - Most NumPy not allowed
  - Code ends up looking more like CUDA than Python

# Testing eigh in CuPy



- We wanted to see how the `np.linalg.eigh` function scaled on an NVIDIA Volta GPU
- Eventually faster than the Cori CPUs, but only at large matrix size

## CPU

```
import numpy as np  
  
cpu_result =  
np.linalg.eigh(data)
```

## GPU

```
import cupy as cp  
  
#move data to gpu  
gpu_data = cp.asarray(data)  
  
gpu_result =  
cp.linalg.eigh(gpu_data)  
  
#move data to cpu  
cpu_result =  
cp.asnumpy(gpu_result)
```

